

results

May 10, 2024

1 Results

```
[ ]: import os
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import scipy.stats as stats
import sys

sys.path.insert(1, "/Users/simon/Documents/II/Dissertation/")
from src.evaluate import get_results_df

%load_ext autoreload
%autoreload 2
```

Collect results

```
[ ]: models = ["Linear", "ARIMA", "RandomForest", "CNN", "LSTM", "ConvLSTM"]
stocks = ["NVDA", "JPM", "HD", "UNH"]

dfs = []
for m in models:
    for s in stocks:
        dfs.append(get_results_df(f"{m}_{s}"))
dfs = pd.concat(dfs)
dfs.to_csv("./results.csv")
```

Loading Linear_NVDA.

Rank 1: trial no. 0, value: 45.0199203187251. Run completed at 2024-04-29
16:52:25.570862

Loading Linear_JPM.

Rank 1: trial no. 0, value: 46.613545816733065. Run completed at 2024-04-29
16:52:27.015066

Loading Linear_HD.

Rank 1: trial no. 0, value: 52.589641434262944. Run completed at 2024-04-29
16:52:28.379346

Loading Linear_UNH.

Rank 1: trial no. 0, value: 47.808764940239044. Run completed at 2024-04-29 16:52:29.689755
Loading ARIMA_NVDA.
Rank 1: trial no. 18, value: 56.97211155378486. Run completed at 2024-04-29 17:11:04.256046
Loading ARIMA_JPM.
Rank 1: trial no. 1, value: 52.98804780876494. Run completed at 2024-04-29 17:15:08.552541
Loading ARIMA_HD.
Rank 1: trial no. 4, value: 52.98804780876494. Run completed at 2024-04-29 17:27:20.232495
Loading ARIMA_UNH.
Rank 1: trial no. 1, value: 47.01195219123506. Run completed at 2024-04-29 17:35:18.839560
Loading RandomForest_NVDA.
Rank 1: trial no. 1, value: 50.59760956175299. Run completed at 2024-04-29 17:44:32.475530
Loading RandomForest_JPM.
Rank 1: trial no. 8, value: 53.38645418326693. Run completed at 2024-04-29 17:51:43.110122
Loading RandomForest_HD.
Rank 1: trial no. 6, value: 54.18326693227091. Run completed at 2024-04-29 17:56:56.054294
Loading RandomForest_UNH.
Rank 1: trial no. 3, value: 55.77689243027888. Run completed at 2024-04-29 17:59:37.088842
Loading CNN_NVDA.
Rank 1: trial no. 11, value: 0.5737051963806152. Run completed at 2024-04-29 20:22:15.814904
Loading CNN_JPM.
Rank 1: trial no. 9, value: 0.5737051963806152. Run completed at 2024-04-29 18:34:47.643948
Loading CNN_HD.
Rank 1: trial no. 0, value: 0.5498008131980896. Run completed at 2024-04-29 18:35:17.552823
Loading CNN_UNH.
Rank 1: trial no. 8, value: 0.5816733241081238. Run completed at 2024-04-29 18:44:16.965952
Loading LSTM_NVDA.
Rank 1: trial no. 0, value: 0.5537848472595215. Run completed at 2024-04-29 18:01:30.993116
Loading LSTM_JPM.
Rank 1: trial no. 3, value: 0.6175298690795898. Run completed at 2024-04-29 18:13:38.976207
Loading LSTM_HD.
Rank 1: trial no. 4, value: 0.6055777072906494. Run completed at 2024-04-29 18:19:41.955427
Loading LSTM_UNH.

Rank 1: trial no. 11, value: 0.5896414518356323. Run completed at 2024-04-29 20:23:47.878337

Loading ConvLSTM_NVDA.

Rank 1: trial no. 2, value: 0.518652081489563. Run completed at 2024-04-29 18:47:06.218729

Loading ConvLSTM_JPM.

Rank 1: trial no. 26, value: 0.5697211027145386. Run completed at 2024-04-29 20:35:01.199259

Loading ConvLSTM_HD.

Rank 1: trial no. 6, value: 0.5231999158859253. Run completed at 2024-04-29 19:04:03.502803

Loading ConvLSTM_UNH.

Rank 1: trial no. 0, value: 0.5577689409255981. Run completed at 2024-04-29 16:26:05.723444

Load results if already exists

```
[ ]: path = "./results.csv"
      if os.path.exists(path):
          dfs = pd.read_csv(path, header=[0, 1], index_col=0)
          dfs["Model Type"] = dfs.index.str.split("_").str[0]
          dfs["Stock"] = dfs.index.str.split("_").str[1]
```

Aggregating by stock

```
[ ]: df = dfs.copy()
      df.drop(columns=["Hyperparameters", "Model Type"]).groupby("Stock").mean().loc[
          ["NVDA", "JPM", "HD", "UNH"]
      ]
```

/var/folders/d7/ktx3dym91yjgj_gpmnfs0rh00000gn/T/ipykernel_22281/248852036.py:2:
PerformanceWarning: dropping on a non-lexsorted multi-index without a level
parameter may impact performance.

```
df.drop(columns=["Hyperparameters", "Model Type"]).groupby("Stock").mean().loc[["NVDA", "JPM", "HD", "UNH"]]
```

```
[ ]: Validation set \
      R2      MSE      RMSE      MAE      p      Accuracy
Stock
NVDA      -0.24753428 0.00196531 0.04376697 0.03503818 0.03884602 49.53519256
JPM        -0.26228733 0.00044194 0.02081327 0.01618928 0.02068996 48.80478088
HD          -0.13828215 0.00044386 0.02092433 0.01598344 0.07250368 51.92563081
UNH         -0.17845246 0.00027840 0.01650261 0.01281574 0.03330397 51.99203187
```

```
Test set \
      Avg. daily return Std. daily return Risk adj. return      R2
Stock
NVDA      -0.00210417      0.03325528      -0.06049100 -0.25203186
```

JPM	-0.00034623	0.01426815	-0.02285721	-0.29839481
HD	-0.00023387	0.01587887	-0.01208317	-0.16416481
UNH	0.00002149	0.01246166	-0.00020706	-0.15088191

	MSE	RMSE	MAE	p	Accuracy
Stock					
NVDA	0.00108453	0.03257765	0.02361088	-0.00395278	52.93333333
JPM	0.00022075	0.01467125	0.01074293	-0.03470588	49.73333333
HD	0.00021324	0.01449804	0.01078709	0.01850120	50.20000000
UNH	0.00020444	0.01417508	0.00998843	0.03085183	51.20000000

	Avg. daily return	Std. daily return	Risk adj. return
Stock			
NVDA	0.00344616	0.02520458	0.13110339
JPM	0.00054191	0.01030267	0.05189790
HD	0.00023232	0.01007062	0.02157074
UNH	0.00015475	0.00913768	0.01866290

```
[ ]: # Plotting
fig, axs = plt.subplots(nrows=1, ncols=2, figsize=(10, 3))

for i, set in enumerate(["Validation set", "Test set"]):
    df = dfs.copy()
    df = df.pivot(index="Stock", columns="Model Type", values=(set, "Accuracy"))
    df = df.loc[["NVDA", "JPM", "HD", "UNH"]] # Re-order rows
    df = df[
        ["Linear", "ARIMA", "RandomForest", "CNN", "LSTM", "ConvLSTM"]
    ] # Re-order columns
    display(df)
    df.plot(kind="bar", ax=axs[i])
    axs[i].set_title(f"{set}")
    axs[i].set_ylabel("Accuracy")
    axs[i].set_xlabel("Stock")
    axs[i].tick_params(labelrotation=0)
    axs[i].get_legend().remove()
    axs[i].set_ylim([30, 60])

axs[1].legend(title="Model Type", bbox_to_anchor=(1.05, 1), loc="upper left")
fig.tight_layout()
plt.show()
```


Model Type	Linear	ARIMA	RandomForest	CNN	LSTM
Stock					
NVDA	45.01992032	56.97211155	50.59760956	48.20717131	48.20717131

JPM	46.61354582	52.98804781	50.99601594	47.41035857	47.41035857
HD	52.58964143	52.98804781	54.18326693	54.58167331	48.60557769
UNH	47.80876494	47.01195219	50.99601594	55.77689243	57.76892430

Model Type	ConvLSTM
Stock	
NVDA	48.20717131
JPM	47.41035857
HD	48.60557769
UNH	52.58964143

Model Type	Linear	ARIMA	RandomForest	CNN	LSTM \
Stock					
NVDA	48.80000000	48.80000000	50.00000000	56.80000000	56.40000000
JPM	50.80000000	48.80000000	46.00000000	46.00000000	50.00000000
HD	46.40000000	46.80000000	53.20000000	47.60000000	53.60000000
UNH	50.40000000	50.80000000	52.00000000	50.40000000	50.80000000

Model Type	ConvLSTM
Stock	
NVDA	56.80000000
JPM	56.80000000
HD	53.60000000
UNH	52.80000000



First, we observe how each model performs on each stock. In terms of accuracy, there is no clear highest-performing model on the validation set. On the other hand, the test set shows the ConvLSTM model outperforming other model types on all stocks.

```
[ ]: # Aggregate by Model Type
df = dfs.copy()
df.drop(columns=["Hyperparameters", "Stock"]).groupby("Model Type").mean().loc[
    ["Linear", "ARIMA", "RandomForest", "CNN", "LSTM", "ConvLSTM"]
]
```

/var/folders/d7/ktx3dym91yjgj_gpmnfs0rh00000gn/T/ipykernel_22281/445887624.py:3:

PerformanceWarning: dropping on a non-lexsorted multi-index without a level parameter may impact performance.

```
df.drop(columns=["Hyperparameters", "Stock"]).groupby("Model
Type").mean().loc[["Linear", "ARIMA", "RandomForest", "CNN", "LSTM",
"ConvLSTM"]]
```

```
[ ]: Validation set
      R2      MSE      RMSE      MAE      p
Model Type
Linear      -0.02066762  0.00065106  0.02361975  0.01858800  0.00687075
ARIMA      -1.02644547  0.00135590  0.03361967  0.02663852  0.03078746
RandomForest -0.03631573  0.00068872  0.02402783  0.01888253  0.06086197
CNN        -0.12797410  0.00068466  0.02448429  0.01912300  0.00949959
LSTM       -0.00628597  0.00064480  0.02347886  0.01832636  0.07087916
ConvLSTM    -0.02214544  0.00066911  0.02378037  0.01848158  0.06177991
```

```
Accuracy Avg. daily return Std. daily return Risk adj. return
Model Type
Linear      48.00796813      -0.00087223      0.01850491      -0.03379087
ARIMA      52.49003984      -0.00030584      0.01509480      -0.02210786
RandomForest 51.69322709      -0.00010573      0.01629282      0.00525725
CNN        51.49402390      -0.00092618      0.01982590      -0.03456280
LSTM       50.49800797      -0.00063349      0.02131060      -0.01753170
ConvLSTM    49.20318725      -0.00115072      0.02276692      -0.04072167
```

```
Test set
      R2      MSE      RMSE      MAE      p
Model Type
Linear      -0.01909141  0.00035809  0.01752181  0.01252915  0.02385229
ARIMA      -1.01053333  0.00071815  0.02469201  0.01912823 -0.00055633
RandomForest -0.10110938  0.00042627  0.01860393  0.01317114 -0.00732916
CNN        -0.16178257  0.00038139  0.01834945  0.01319566 -0.04282749
LSTM       -0.00462845  0.00035063  0.01737026  0.01233547  0.02478478
ConvLSTM    -0.00106495  0.00034991  0.01734557  0.01233434  0.00839879
```

```
Accuracy Avg. daily return Std. daily return Risk adj. return
Model Type
Linear      49.10000000      0.00090831      0.01171694      0.05495246
ARIMA      48.80000000      0.00058564      0.01129133      0.04494016
RandomForest 50.30000000      0.00086942      0.01082230      0.07098626
CNN        50.20000000      0.00119086      0.01503570      0.03647862
LSTM       52.70000000      0.00125459      0.01608904      0.04510046
ConvLSTM    55.00000000      0.00175389      0.01711801      0.08239443
```


Looking at the mean accuracies of each model for the test set, we observe that the hybrid ConvLSTM is improved across all metrics over its CNN and LSTM predecessors. The second-best model

is the LSTM model, followed by the CNN and Random Forest models. The ConvLSTM model, relative to the other models performs much better. Now turning our attention to the standard deviation of daily return, we observe that the statistical models all similarly have lower variation than the deep learning models. This suggests the trading decisions derived from the deep learning models result in riskier portfolios, whose values vary more than those of the statistical models. However, this is counter-balanced by the average daily return being higher for the deep learning models. Consequently, all models have similar risk adjusted returns, except the ConvLSTM model which is notably higher. Similar rankings can be observed from the traditional machine learning metrics, with higher accuracy correlating to higher average daily return, but as noted in Section X, this is not always true.

```
[ ]: # Plotting
fig, axs = plt.subplots(nrows=1, ncols=2, figsize=(10, 3))

for i, set in enumerate(["Validation set", "Test set"]):
    df = dfs.copy()
    df = df.pivot(
        index="Stock", columns="Model Type", values=(set, "Avg. daily return")
    )
    df = df.loc[["NVDA", "JPM", "HD", "UNH"]] # Re-order rows
    df = df[
        ["Linear", "ARIMA", "RandomForest", "CNN", "LSTM", "ConvLSTM"]
    ] # Re-order columns
    df.plot(kind="bar", ax=axs[i])
    axs[i].set_title(f"{set}")
    axs[i].set_ylabel("Avg. daily return")
    axs[i].set_xlabel("Stock")
    axs[i].tick_params(labelrotation=0)
    axs[i].get_legend().remove()

axs[1].legend(title="Model Type", bbox_to_anchor=(1.05, 1), loc="upper left")
fig.tight_layout()
plt.show()
```



Viewing the average daily return for each of the models across the selected stocks shows large

variation in NVIDIA stock. This can be explained with the high volatility of the stock during the validation and test set periods (2023 and 2024, respectively), where many external factors and events influenced sharp jumps in price, as can be seen in Figure X.

In summary, comparing the models to each other, we conclude that ConvLSTM model is the best, whereas the linear regression and ARIMA models are the worst. Now we shift our attention to comparing the models to a baseline random walk model.

```
[ ]: import seaborn as sns
      from src.evaluate import random_walk

      stocks = ["NVDA", "JPM", "HD", "UNH"]

      random_df = []
      for s in stocks:
          random_df.append(random_walk(s))
      random_df = pd.concat(random_df)
      random_df
```

```
[ ]:      R2      MSE      RMSE      MAE      p      Accuracy \
0      -1.05775320 0.00178245 0.04221912 0.03158757 -0.01961993 50.00000000
1      -1.08001233 0.00180174 0.04244685 0.03182528 0.00067114 49.60000000
2      -0.95982756 0.00169763 0.04120230 0.03102806 0.06319344 53.60000000
3      -0.93820740 0.00167890 0.04097440 0.03150312 0.02393293 48.40000000
4      -0.85960718 0.00161082 0.04013499 0.03134208 0.12671680 53.20000000
..      ...      ...      ...      ...      ...      ...
995     -1.87922097 0.00051145 0.02261517 0.01745124 -0.03715600 52.80000000
996     -1.87011769 0.00050983 0.02257939 0.01814187 0.14378865 55.20000000
997     -2.25343121 0.00057792 0.02403992 0.01892841 -0.03550563 50.40000000
998     -1.73058412 0.00048504 0.02202369 0.01778347 0.07138073 48.40000000
999     -2.14836870 0.00055926 0.02364858 0.01857498 -0.05694255 46.40000000
```

	Avg. daily return	Std. daily return	Risk adj. return	Stock
0	0.00259022	0.02522597	0.10268063	NVDA
1	0.00243437	0.02148658	0.11329744	NVDA
2	0.00316054	0.01762301	0.17934184	NVDA
3	0.00350995	0.02116178	0.16586284	NVDA
4	0.00413856	0.02482279	0.16672422	NVDA
..
995	0.00010428	0.00993493	0.01049597	UNH
996	0.00093807	0.00932627	0.10058320	UNH
997	0.00012910	0.00897467	0.01438475	UNH
998	0.00031571	0.00986056	0.03201728	UNH
999	-0.00011882	0.00851556	-0.01395273	UNH


[4000 rows x 10 columns]


```
[ ]: df = random_df
display(df)
fig, axs = plt.subplots(nrows=1, ncols=4, sharey=True, figsize=(12, 3))
sns.violinplot(data=df, x="RMSE", y="Stock", ax=axs[0], linewidth=1)
sns.violinplot(data=df, x="Accuracy", y="Stock", ax=axs[1], linewidth=1).set(
    ylabel=None
)
sns.violinplot(data=df, x="Avg. daily return", y="Stock", ax=axs[2],
    linewidth=1).set(
    ylabel=None
)
sns.violinplot(data=df, x="Risk adj. return", y="Stock", ax=axs[3],
    linewidth=1).set(
    ylabel=None
)
fig.tight_layout()
```

	R2	MSE	RMSE	MAE	p	Accuracy \
0	-1.05775320	0.00178245	0.04221912	0.03158757	-0.01961993	50.00000000
1	-1.08001233	0.00180174	0.04244685	0.03182528	0.00067114	49.60000000
2	-0.95982756	0.00169763	0.04120230	0.03102806	0.06319344	53.60000000
3	-0.93820740	0.00167890	0.04097440	0.03150312	0.02393293	48.40000000
4	-0.85960718	0.00161082	0.04013499	0.03134208	0.12671680	53.20000000
..
995	-1.87922097	0.00051145	0.02261517	0.01745124	-0.03715600	52.80000000
996	-1.87011769	0.00050983	0.02257939	0.01814187	0.14378865	55.20000000
997	-2.25343121	0.00057792	0.02403992	0.01892841	-0.03550563	50.40000000
998	-1.73058412	0.00048504	0.02202369	0.01778347	0.07138073	48.40000000
999	-2.14836870	0.00055926	0.02364858	0.01857498	-0.05694255	46.40000000

	Avg. daily return	Std. daily return	Risk adj. return	Stock
0	0.00259022	0.02522597	0.10268063	NVDA
1	0.00243437	0.02148658	0.11329744	NVDA
2	0.00316054	0.01762301	0.17934184	NVDA
3	0.00350995	0.02116178	0.16586284	NVDA
4	0.00413856	0.02482279	0.16672422	NVDA
..
995	0.00010428	0.00993493	0.01049597	UNH
996	0.00093807	0.00932627	0.10058320	UNH
997	0.00012910	0.00897467	0.01438475	UNH
998	0.00031571	0.00986056	0.03201728	UNH
999	-0.00011882	0.00851556	-0.01395273	UNH

[4000 rows x 10 columns]



```
[ ]: random_df.groupby("Stock").mean()
```

```
[ ]:
```

	R2	MSE	RMSE	MAE	p	Accuracy \
Stock						
HD	-1.46025088	0.00045065	0.02120991	0.01665570	-0.00154551	50.02640000
JPM	-3.20422480	0.00071479	0.02671067	0.02112854	-0.00347414	49.91320000
NVDA	-1.08580890	0.00180676	0.04247356	0.03254093	0.00374187	50.04440000
UNH	-2.17078680	0.00056324	0.02371141	0.01861061	0.00136204	50.04960000

	Avg. daily return	Std. daily return	Risk adj. return
Stock			
HD	0.00023891	0.00968711	0.02529536
JPM	0.00053071	0.00925316	0.05734307
NVDA	0.00254700	0.02185612	0.11533465
UNH	0.00013011	0.00954605	0.01351766

We can see the evaluation metric distributions of the random walk model across the selected stocks. Notably, the RMSE for a random walk model is much higher for NVIDIA, which reflects the high price volatility and unpredictability. However, this also meant a higher return during the test period of NVIDIA stock price. The accuracies have a mean of 50 may initially seem like a discrepancy as daily returns differ, but this is simply due to the compounding nature of stock returns.

```
[ ]: df = random_df.drop(columns=["Stock"]).groupby(random_df.index).mean()
display(df)

fig, axs = plt.subplots(nrows=1, ncols=4, figsize=(12, 3))
for i, metric in enumerate(
    ["RMSE", "Accuracy", "Avg. daily return", "Risk adj. return"]
):
    sns.histplot(df, x=metric, kde=True, ax=axs[i], stat="density")
    xmin, xmax = axs[i].get_xlim()
    x = np.linspace(xmin, xmax, 100)

    mean = df[metric].mean()
    std = df[metric].std()
    p = stats.norm.pdf(x, mean, std)
```

```

    axs[i].plot(x, p, "k", linewidth=1, label="Normal Distribution",
↳linestyle="dashed")

    if metric == "RMSE":
        ci = stats.norm.ppf(0.05, mean, std)
        axs[i].fill_between(x, p, where=(x <= ci), color="red", alpha=0.5)
    else:
        ci = stats.norm.ppf(0.95, mean, std)
        axs[i].fill_between(x, p, where=(x >= ci), color="red", alpha=0.5)
    axs[i].axvline(ci, color="red", linestyle="--", label=f"x = {ci:.4f}")
    axs[i].legend(loc="upper left")

fig.tight_layout()

```

	R2	MSE	RMSE	MAE	p	Accuracy \
0	-1.95546575	0.00087338	0.02828586	0.02198506	0.03776665	50.20000000
1	-2.10949358	0.00090515	0.02895658	0.02238151	-0.02755672	50.20000000
2	-1.89863779	0.00084762	0.02801149	0.02169403	0.04287293	53.70000000
3	-1.99163340	0.00086011	0.02832100	0.02233371	-0.02568557	48.50000000
4	-1.95365457	0.00083973	0.02802742	0.02206717	0.03585953	52.00000000
..
995	-2.06676229	0.00091894	0.02894076	0.02224064	-0.03988918	49.10000000
996	-1.95490660	0.00091105	0.02877105	0.02234196	0.00612882	49.40000000
997	-2.08609802	0.00089164	0.02877677	0.02272362	0.01443309	49.70000000
998	-1.87897104	0.00088667	0.02838867	0.02225543	0.00235307	48.70000000
999	-2.07871308	0.00092113	0.02905704	0.02281539	-0.02309645	48.30000000

	Avg. daily return	Std. daily return	Risk adj. return
0	0.00110122	0.01366336	0.07209553
1	0.00079224	0.01182458	0.04968050
2	0.00127771	0.01142796	0.10029224
3	0.00093262	0.01240810	0.04620775
4	0.00133348	0.01324890	0.07560883
..
995	0.00075601	0.01311611	0.03696599
996	0.00087363	0.01305440	0.05726713
997	0.00105711	0.01168628	0.07901336
998	0.00079473	0.01331819	0.02982082
999	0.00040858	0.01092924	0.02214841

[1000 rows x 9 columns]

```

/Users/simon/anaconda3/envs/proj/lib/python3.9/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.

```

```

    with pd.option_context('mode.use_inf_as_na', True):
/Users/simon/anaconda3/envs/proj/lib/python3.9/site-

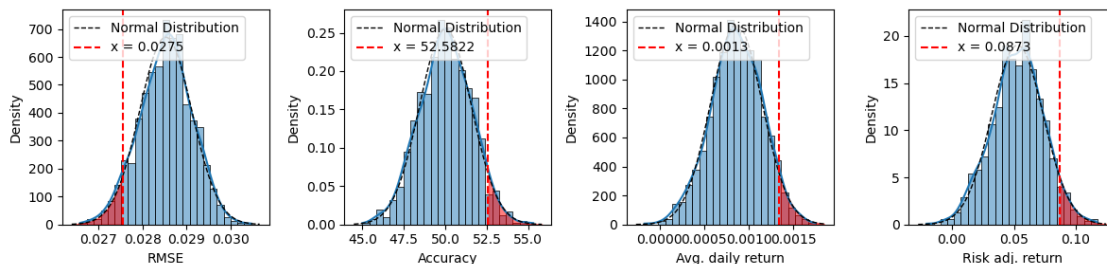
```

packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
/Users/simon/anaconda3/envs/proj/lib/python3.9/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
/Users/simon/anaconda3/envs/proj/lib/python3.9/site-
packages/seaborn/_oldcore.py:1119: FutureWarning: use_inf_as_na option is
deprecated and will be removed in a future version. Convert inf values to NaN
before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```



```
[ ]: df = dfs.copy()
df.drop(columns=["Hyperparameters", "Stock"]).groupby("Model Type").mean().loc[
    ["Linear", "ARIMA", "RandomForest", "CNN", "LSTM", "ConvLSTM"]
][ "Test set"]
```

/var/folders/d7/ktx3dym91yjgj_gpmnfs0rh00000gn/T/ipykernel_58740/377110121.py:2: PerformanceWarning: dropping on a non-lexsorted multi-index without a level parameter may impact performance.

```
df.drop(columns=["Hyperparameters", "Stock"]).groupby("Model
Type").mean().loc[["Linear", "ARIMA", "RandomForest", "CNN", "LSTM",
"ConvLSTM"]][ "Test set"]
```

```
[ ]:
R2      MSE      RMSE      MAE      p \
Model Type
Linear   -0.01909141  0.00035809  0.01752181  0.01252915  0.02385229
ARIMA    -1.01053333  0.00071815  0.02469201  0.01912823 -0.00055633
RandomForest -0.10110938  0.00042627  0.01860393  0.01317114 -0.00732916
CNN      -0.16178257  0.00038139  0.01834945  0.01319566 -0.04282749
LSTM     -0.00462845  0.00035063  0.01737026  0.01233547  0.02478478
ConvLSTM -0.00106495  0.00034991  0.01734557  0.01233434  0.00839879
```

```
Accuracy Avg. daily return Std. daily return \
```

Model Type			
Linear	49.10000000	0.00090831	0.01171694
ARIMA	48.80000000	0.00058564	0.01129133
RandomForest	50.30000000	0.00086942	0.01082230
CNN	50.20000000	0.00119086	0.01503570
LSTM	52.70000000	0.00125459	0.01608904
ConvLSTM	55.00000000	0.00175389	0.01711801

Risk adj. return	
Model Type	
Linear	0.05495246
ARIMA	0.04494016
RandomForest	0.07098626
CNN	0.03647862
LSTM	0.04510046
ConvLSTM	0.08239443

Performing the statistical tests, we see that only the ConvLSTM model is statistically better than the random walk model in terms of RMSE, accuracy and average daily return. However, the ConvLSTM model does not fall in the critical region for risk adjusted return. This suggests that although the ConvLSTM is able to generate higher returns, it comes with more risk. No other models are statistically significant across average daily return, although the LSTM model also outperforms the random walk model in terms of accuracy. Finally, all models have an improved RMSE, which supports the use of financial metrics, demonstrating that a good RMSE does not always translate to good trading performance.

1.1 Risk analysis

```
[ ]: df = dfs.copy()
df = (
    df.drop(columns=["Stock", "Hyperparameters"])
    .groupby("Model Type")
    .mean()["Test set"]
)
display(df[["Avg. daily return", "Std. daily return"]])

fig, ax = plt.subplots(figsize=(4, 4))
ax.scatter(x=df["Avg. daily return"], y=df["Std. daily return"])
plt.xlabel("Avg. daily return")
plt.ylabel("Risk (std. of daily return)")
ax.grid(True, alpha=0.2)

for label, x, y in zip(df.index, df["Avg. daily return"], df["Std. daily_
    ↪return"]):
    plt.annotate(
        label,
        xy=(x, y),
```

```

        xytext=(10, 0),
        textcoords="offset points",
        ha="right",
        va="bottom",
    )

```

```

/var/folders/d7/ktx3dym91yjgj_gpmnfs0rh00000gn/T/ipykernel_22281/2130644621.py:2
: PerformanceWarning: dropping on a non-lexsorted multi-index without a level
parameter may impact performance.

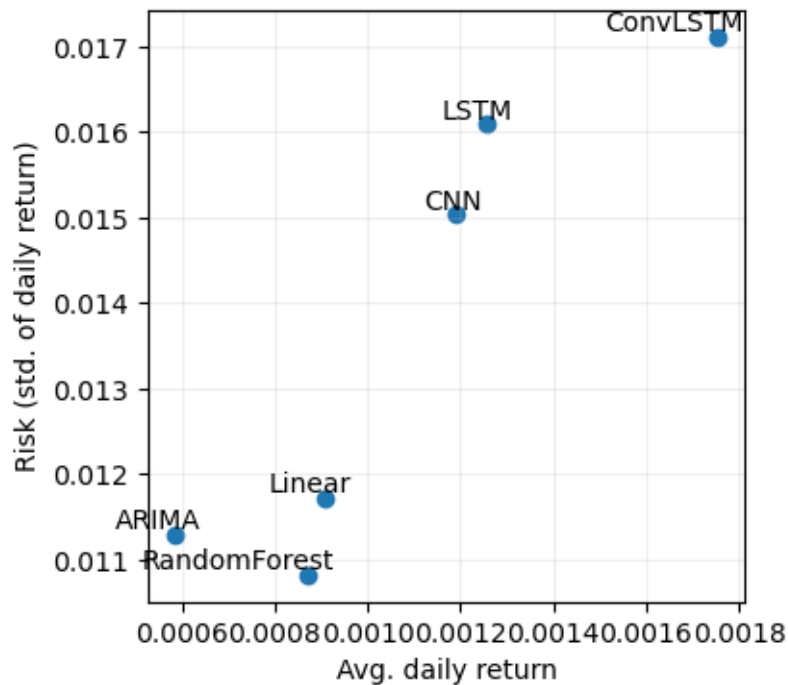
```

```

df = df.drop(columns=["Stock", "Hyperparameters"]).groupby("Model
Type").mean()["Test set"]

```

Model Type	Avg. daily return	Std. daily return
ARIMA	0.00058564	0.01129133
CNN	0.00119086	0.01503570
ConvLSTM	0.00175389	0.01711801
LSTM	0.00125459	0.01608904
Linear	0.00090831	0.01171694
RandomForest	0.00086942	0.01082230



```

[ ]: df = dfs.copy()
df = (
    df.drop(columns=["Model Type", "Hyperparameters"])
    .groupby("Stock")

```

```

        .mean()["Test set"]
    )
display(df[["Avg. daily return", "Std. daily return"]])

fig, ax = plt.subplots(figsize=(4, 4))
ax.scatter(x=df["Avg. daily return"], y=df["Std. daily return"])
plt.xlabel("Avg. daily return")
plt.ylabel("Risk (std. of daily return)")
ax.grid(True, alpha=0.2)
ax.set_ylim((0.007, 0.03))
ax.set_xlim((-0.0005, 0.0041))

for label, x, y in zip(df.index, df["Avg. daily return"], df["Std. daily_
↪return"]):
    plt.annotate(
        label,
        xy=(x, y),
        xytext=(0, 0),
        textcoords="offset points",
        ha="right",
        va="bottom",
    )

```

/var/folders/d7/ktx3dym91yjgj_gpmnfs0rh00000gn/T/ipykernel_22281/1779348945.py:2
: PerformanceWarning: dropping on a non-lexsorted multi-index without a level
parameter may impact performance.

```

df = df.drop(columns=["Model Type",
"Hyperparameters"]).groupby("Stock").mean()["Test set"]

```

	Avg. daily return	Std. daily return
Stock		
HD	0.00023232	0.01007062
JPM	0.00054191	0.01030267
NVDA	0.00344616	0.02520458
UNH	0.00015475	0.00913768

